


N.A.T. GmbH

**TCP/IP – Socket Library
Programmer's Reference
Version 1.6**




**N.A.T. GmbH
Kamillenweg 22
D-53757 Sankt Augustin**

Phone: ++49/2241/3989-0

Fax: ++49/2241/3989-10

E-Mail: support@nateurope.com

Internet: <http://www.nateurope.com>



Disclaimer

The following documentation, compiled by N.A.T. GmbH (henceforth called N.A.T.), represents the current status of the products development. The documentation is updated on a regular basis. Any changes which might ensue, including those necessitated by updated specifications, are considered in the latest version of this documentation. N.A.T. is under no obligation to notify any person, organisation, or institution of such changes or to make these changes public in any other way.

We must caution you, that this publication could include technical inaccuracies or typographical errors.

N.A.T. offers no warranty, either expressed or implied, for the contents of this documentation or for the product described therein, including but not limited to the warranties of merchantability or the fitness of the product for any specific purpose.

In no event, will N.A.T. be liable for any loss of data or for errors in data utilisation or processing resulting from the use of this product or the documentation. In particular, N.A.T. will not be responsible for any direct or indirect damages (including lost profits, lost savings, delays or interruptions in the flow of business activities, including but not limited to, special, incidental, consequential, or other similar damages) arising out of the use of or inability to use this product or the associated documentation, even if N.A.T. or any authorised N.A.T. representative has been advised of the possibility of such damages.

The use of registered names, trademarks, etc. in this publication does not imply, even in the absence of a specific statement, that such names are exempt from the relevant protective laws and regulations (patent laws, trade mark laws, etc.) and therefore free for general use. In no case does N.A.T. guarantee that the information given in this documentation is free of such third-party rights.

Neither this documentation nor any part thereof may be copied, translated, or reduced to any electronic medium or machine form without the prior written consent from N.A.T. GmbH.

This product (and the associated documentation) is governed by the N.A.T. General Conditions and Terms of Delivery and Payment.






Table of Contents

1	INTRODUCTION	6
2	SOCKET INTERFACE	7
2.1	GENERAL DIFFERENCE BETWEEN TCP/IP AND ISO/OSI	8
2.2	ACCEPT()	9
2.3	BIND()	11
2.4	CLOSE()	13
2.5	CONNECT()	14
2.6	GETHOSTNAME()	16
2.7	GETPEERNAME()	18
2.8	GETSOCKNAME()	19
2.9	GETSOCKOPT()	20
2.10	IOCTL()	23
2.11	LISTEN()	24
2.12	READ()	25
2.13	RECV()	26
2.14	RECVFROM()	28
2.15	RECVMSG()	30
2.16	SELECT	32
2.17	SEND()	34
2.18	SENDX()	36
2.19	SENDMSG()	37
	MSG_OOB	38
2.20	SENDTO()	40
2.21	SHUTDOWN()	42
2.22	SETSOCKOPT()	43
2.23	SOCKET()	48
2.24	SYSREQMEM()	51
2.25	SYSRETMEM()	52
2.26	WRITE()	53
3	CONVERSION ROUTINES	54
3.1	HTONL()	55
3.2	HTONS()	56
3.3	NTOHL()	57
3.4	NTOHS()	58
4	ADDRESS MANIPULATION FUNCTIONS	59
4.1	INTERNET ADDRESSES IN DOT NOTATION	60
4.2	INET_ADDR()	61
4.3	INET_LNAOF()	62
4.4	INET_MAKEADDR()	63
4.5	INET_NETOF()	64
4.6	INET_NETWORK()	65
4.7	INET_NTOA()	66




5	NETWORK DATA BASIS	67
5.1	HOSTS	69
5.1.1	<i>The struct hostent structure</i>	70
5.1.2	<i>gethostbyaddr()</i>	71
5.1.3	<i>gethostbyname()</i>	72
5.1.4	<i>gethostent()</i>	74
5.1.5	<i>sethostent()</i>	75
5.1.5	<i>endhostent()</i>	76
5.2	NETWORKS	77
5.2.1	<i>The struct netent structure</i>	78
5.2.2	<i>getnetbyaddr()</i>	79
5.2.2	<i>getnetbyname()</i>	80
5.2.4	<i>getnetent()</i>	81
5.2.4	<i>setnetent()</i>	82
5.2.6	<i>endnetent()</i>	83
5.3	PROTOCOLS	84
5.3.1	<i>The struct protoent structure</i>	85
5.3.2	<i>getprotobyname()</i>	86
5.3.3	<i>getprotobynumber()</i>	87
5.3.4	<i>getprotoent()</i>	88
5.3.5	<i>setprotoent()</i>	89
5.3.6	<i>endprotoent()</i>	90
5.4	SERVICES	91
5.4.1	<i>The struct servent structure</i>	93
5.4.2	<i>getservbyname()</i>	94
5.4.3	<i>getservbyport()</i>	95
5.4.4	<i>getservent()</i>	96
5.4.5	<i>setservent()</i>	97
5.4.6	<i>endservent()</i>	98



1 Introduction

This manual provides a description of the programmer's interface to the N.A.T. TCP/IP and N.A.T. OSI/TP protocol software. This includes the socket interface, conversion and address manipulating routines and functions concerning the network data basis.

The interface can be used on both the TCP/IP and the ISO/OSI protocol stack. There are only slight differences as to the implementation on these different protocols. In case the usage of a specific function should be different on TCP/IP and ISO/OSI, the differences will be explained; any other functions can be used in the same way on both protocols.




2 Socket Interface

The following chapter presents in detail the individual functions of the socket interface in detail. These functions may be found in the library socklib.1. The individual functions are:

accept()	Accept a request for a socket and create it.
bind()	Bind a name to a socket.
close()	Close a connection to a socket
connect()	Initiate a connection to a socket
gethostname()	Return the name of your own host system
getpeername()	Return the socket names of the connected systems
getsockname()	Return the name of your own socket
getsockopt()	View the values of a socket's options
ioctl()	Manipulate the device parameters
listen()	Take requests for a connection to a socket
read()	Read data from a socket
recv()	Read data from a socket
recvfrom()	Read data from a socket
recvmsg()	Read data from a socket
select()	Simultaneously supervise multiple I/O paths
send()	Write data to a socket
sendx()	extended send() call (only on FDDI29/ETH29)
sendmsg()	Write data to a socket
sendto()	Write data to a socket
setsockopt()	Set socket options
shutdown()	Shutdown a portion of a full-duplex connection
socket()	Create a communications end point – socket
SysReqMem()	allocate system memory (only on FDDI29/ETH29)
SysRetMem()	return system memory (only on FDDI29/ETH29)
write()	Write data to a socket

The functions `gethostname()`, `ioctl()`, `read()`, `select()`, `sendx()`, `SysReqMem()`, `SysRetMem()` and `write()` are extensions to the standard BSD socket interface.



2.1 General difference between TCP/IP and ISO/OSI

Between the implementations of the socket interface on TCP/IP and ISO/OSI there is one difference which is of general nature: wherever the struct `sockaddr` structure appears in the description of the socket interface in this manual (as it is used with TCP/IP) it has to be replaced by the struct `sockaddr_iso` structure when working with the OSI socket interface. This has the simple reason that the syntax of OSI addresses and TCP/IP addresses is different

The struct `sockaddr` structure used on TCP/IP is defined as follows:

```
struct sockaddr_in {
    short sin_family;
    u_short sin_port;
    struct in_addr sin_addr;
    char sin_zero[8];
};
```

This is the struct `in_addr` structure:


```
struct in_addr {
    u_long s_addr;
};
```

The following is the struct `sockaddr_iso` structure used on ISO/OSI:

```
struct sockaddr_iso {
    u_char siso_len;
    u_char siso_family;
    u_char siso_plen;
    u_char siso_slen;
    u_char siso_tlen;
    struct iso_addr siso_addr;
    u_char siso_pad[6];
};
```

This is the struct `iso_addr` structure:

```
struct iso_addr {
    u_char isoa_len;           /* length (in bytes) */
    char isoa_genaddr[20];     /* general opaque address */
};
```

2.2 accept()

NAME

accept() - accept a connection on a socket

SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
ns = accept(s, addr, addrlen)
      int ns, s;
      struct sockaddr *addr;
      int addrlen;
```

DESCRIPTION


The argument *s* is a socket that has been created with `socket()`, bound to an address with `bind()`, and is listening for connections after a `listen()`. `Accept` extracts the first connection on the queue of pending connections, creates a new socket with the same properties of *s* and allocates a new file descriptor, *ns*, for the socket. If no pending connections are present on the queue, and the socket is not marked as non-blocking, `accept` blocks the caller until a connection is present. If the socket is marked non-blocking and no pending connections are present on the queue, `accept` returns an error as described below. The accepted socket, *ns*, may not be used to accept more connections. The original socket *s* remains open.

The argument *addr* is a result parameter that is filled in with the address of the connecting entity, as known to the communications layer. The exact format of the *addr* parameter is determined by the domain in which the communication is occurring. The *addrlen* is a value-result parameter; it should initially contain the amount of space pointed to by *addr*; on return it will contain the actual length (in bytes) of the address returned. This call is used with connection-based socket types, currently with `SOCK_STREAM` for TCP/IP or `SOCK_SEQPACKET` for ISO/OSI.

It is possible to `select()` a socket for the purposes of doing an `accept()` by selecting it for `read()`.

RETURN VALUE

The call returns -1 on error. If it succeeds, it returns a non-negative integer that is a descriptor for the accepted socket.




ERRORS

The call fails if:

[EBADF]	The descriptor is invalid.
[EINVAL]	A listen() must be executed before the accept().
[EWOULDBLOCK]	The socket s is marked non-blocking and no connections are present to be accepted.
[ECONNABORTED]	The socket s is cannot receive any more data.
[EOPNOTSUPP]	The referenced socket is not of type SOCK_STREAM or SOCK_SEQPACKET.
[ENOBUFS]	There are not enough system resources to execute the function.

SEE ALSO

bind(), connect(), listen(), select(), socket()



2.3 bind()

NAME

bind() - bind a name to a socket

SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
bind(s, name, namelen)
    int s;
    struct sockaddr *name;
    int namelen;
```

DESCRIPTION

Bind assigns a name to an unnamed socket. When a socket is created with socket() it exists in a name space (address family) but has no name assigned. Bind requests that name be assigned to the socket s.


RETURN VALUE

If the bind is successful, a 0 value is returned. A return value of -1 indicates an error, which is further specified in the global errno.

ERRORS


The call fails if:

[EBADF]	S is not a valid path number for a socket descriptor.
[EADDRNOTAVAIL]	The specified address (name) is not available from the local machine.
[EADDRINUSE]	The specified address (name) is already in use.
[E_ILLARG]	The specified length namelen does not match the length of names in the specified address family.
[EACCES]	The requested address is protected, and the current user has inadequate permission to access it.
[ENOBUFFS]	There are not enough system resources to execute the function.



SEE ALSO

`connect()`, `listen()`, `socket()`, `getsockname()`



2.4 close()

NAME

close() - delete a descriptor

SYNOPSIS

close(s)
int s;

DESCRIPTION

The close call corresponds to the normal OS-9 close()- function and is used in exactly the same manner. The parameter *s* is a socket that has been created using either socket() or accept(). When a socket is no longer in use, close is used to break the connection and free all the system resources which were used by this socket.

RETURN VALUE

Upon successful completion, a value of 0 is returned. Otherwise, a value of -1 is returned and the global integer variable errno is set to indicate the error.


ERRORS

The call fails if:

[EBADF] *s* is not an active descriptor.

SEE ALSO

accept(), socket()



2.5 connect()

NAME

connect() - initiate a connection on a socket

SYNOPSIS

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int connect(s, name, namelen)
    int s;
    struct sockaddr *name;
    int namelen;
```


DESCRIPTION

The parameter *s* is a socket. If it is of type `SOCK_DGRAM`, then this call specifies the peer with which the socket is to be associated; this address is that to which datagrams are to be sent, and the only address from which datagrams are to be received. If the socket is of type `SOCK_STREAM` (or `SOCK_SEQPACKET` on ISO/OSI), then this call attempts to make a connection to another socket. The other socket is specified by *name*, which is an address in the communications space of the socket.

Each communications space interprets the *name* parameter in its own way. Generally, stream sockets may successfully connect only once; datagram sockets may use connect multiple times to change their association. Datagram sockets may dissolve the association by connecting to an invalid address, such as a null address.

RETURN VALUE

If the connection or binding succeeds, then 0 is returned. Otherwise a -1 is returned, and a more specific error code is stored in `errno`.




ERRORS

The call fails if:

[EBADF]	S is not a valid descriptor.
[EADDRINUSE]	The address is already in use.
[EADDRNOTAVAIL]	The specified address is not available on this machine.
[EAFNOSUPPORT]	Addresses in the specified address family cannot be used with this socket.
[EALREADY]	The socket is non-blocking and a previous connection attempt has not yet been completed.
[ECONNREFUSED]	The attempt to connect was forcefully rejected.
[EINPROGRESS]	The socket is non-blocking and the connection cannot be completed immediately. It is possible to select() for completion by selecting the socket for writing.
[EISCONN]	The socket is already connected.
[ENETUNREACH]	The network isn't reachable from this host.
[ENOBUFS]	There are not enough system resources available to execute the function.
[E_ILLARG]	The specified namelen parameter specifies does not match the length of the names in the specified address family.

SEE ALSO

accept(), select(), socket(), getsockname()



2.6 gethostname()

NAME

gethostname(), sethostname() - get/set name of current host

SYNOPSIS

```
int gethostname(name, namelen)  
    char *name;  
    int namelen;
```

```
int sethostname(name, namelen)  
    char *name;  
    int namelen;
```

DESCRIPTION

Gethostname() returns the standard host name for your own processor. The parameter *namelen* specifies the size of the name array. The returned name is null-terminated unless insufficient space is provided.


Gethostname() searches for the name by first checking the assigned location in the sock-descriptor and then (assuming that there was no string found) in reading the file /DD/ETC/systemid. If an entry for the host name is not found in either location, an ENOENAME error is returned. If an entry for the name is found, the string is copied to the parameter *name*.

The host name is normally automatically written in the file /DD/ETC/systemid during the software configuration by the utility natconfig. If the file does not already exist natconfig will create it. Choosing natconfig's sequence of menu selections Configuration and then Software will result in a dialog requesting the name of your own host processor. (It is also possible use an ASCII text editor to make the entry in the file /DD/ETC/systemid.) The file /DD/ETC/systemid should consist of just the host name and a terminating carriage return <CR>.

If you wish to enter the name in the sock-descriptor, you must use an editor to make the entry on the line labeled HNAME. The descriptor must then be generated again using make. The descriptor's source and make files are located in the subdirectory SOURCE in the main TCPIP directory.

RETURN VALUE

If the call succeeds a value of 0 is returned. If the call fails, then a value of -1 is returned and an error code is placed in the global location *errno*.



ERRORS


The call fails if:

[ENOHNAME]

No host name was found in either the sock descriptor or the file /DD/ETC/systemid.

NOTE

Host names are limited to a length of 64 characters.



2.7 getpeername()

NAME

getpeername() - get name of connected peer

SYNOPSIS

```
int getpeername(s, name, namelen)  
    int s;  
    struct sockaddr *name;  
    int *namelen;
```

DESCRIPTION

Getpeername returns the name of the peer connected to socket *s*. The *namelen* parameter should be initialized to indicate the amount of space pointed to by *name*. On return, it contains the actual size of the name returned (in bytes). The name is truncated if the buffer provided is too small.

RETURN VALUE

A 0 is returned if the call succeeds, -1 if it fails and an appropriate error code is returned in *errno*.


ERRORS

The call fails if:

[EBADF]	The argument <i>s</i> is not a valid descriptor.
[ENOTCONN]	The socket is not connected.
[ENOBUFS]	Insufficient resources were available in the system to perform the operation.

SEE ALSO

accept(), bind(), socket(), getsockname()



2.8 getsockname()

NAME

getsockname() - get socket name

SYNOPSIS

```
int getsockname(s, name, namelen)  
int s;  
struct sockaddr *name;  
int *namelen;
```

DESCRIPTION

Getsockname returns the current name for the specified socket. The namelen parameter should be initialized to indicate the amount of space pointed to by name. On return it contains the actual size of the name returned (in bytes). The name will be truncated, if the specified buffer is too small.

RETURN VALUE

A 0 is returned if the call succeeds, -1 if it fails and an appropriate error code is returned in errno.


ERRORS

The call fails if:

[EBADF]	The argument s is not a valid descriptor.
[ENOBUFS]	Insufficient resources were available in the system to perform the operation.

SEE ALSO

bind(), socket()



2.9 getsockopt()

NAME

getsockopt() - get socket options

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int getsockopt(s, level, optname, optval, optlen)
```

```
    int s, level, optname;
```

```
    char *optval;
```

```
    int *optlen;
```

DESCRIPTION

Getsockopt can be used to view the options associated with a socket. Options may exist at multiple protocol levels; they are always present at the uppermost ``socket" level.

When viewing socket options, the level at which the option resides and the name of the option must be specified. To view options at the ``socket" level, level is specified as SOL_SOCKET. To view options at any other level, the protocol number of the appropriate protocol controlling the option must be entered. For example, to indicate that an option is to be interpreted by the TCP protocol, level should be set to the protocol number of TCP; see getprotoent().


For getsockopt(), the parameters optval and optlen identify a buffer in which the value for the requested option() are to be returned. For getsockopt, optlen is a value-result parameter, initially containing the size of the buffer pointed to by optval, and modified on return to indicate the actual size of the value returned. If no option value is to be supplied or returned, optval may be supplied as 0.

Optname and any specified options are passed uninterpreted to the appropriate protocol module for interpretation.

The include file <socket.h> contains definitions for ``socket" level options, described below. Options at other protocol levels vary in format and name; consult the corresponding entries in the appropriate manuals.

Manual difference:

Most socket-level options take an int parameter for optval. For setsockopt, the parameter should non-zero to enable a boolean option, or zero if the option is to be disabled. SO_LINGER uses a linger structure parameter, defined in <socket.h>, which



specifies the desired state of the option and the linger interval (see below).

The following options are recognized at the socket level. Except as noted, each may be examined with `getsockopt` and set with `setsockopt`.

<code>SO_DEBUG</code>	toggle recording of debugging information
<code>SO_REUSEADDR</code>	toggle local address reuse
<code>SO_KEEPALIVE</code>	toggle keep connections alive
<code>SO_DONTROUTE</code>	toggle routing bypass for outgoing messages
<code>SO_LINGER</code>	linger on close if data present
<code>SO_BROADCAST</code>	toggle permission to transmit broadcast messages
<code>SO_OOBINLINE</code>	toggle reception of out-of-band data in band
<code>SO_SNDBUF</code>	set buffer size for output
<code>SO_RCVBUF</code>	set buffer size for input
<code>SO_TYPE</code>	get the type of the socket (get only)
<code>SO_ERROR</code>	get and clear error on the socket (get only)

`SO_DEBUG` enables debugging in the underlying protocol modules. The data is written to specific address areas and can later be viewed using the program `netprint`.

`SO_REUSEADDR` indicates that the rules used in validating addresses supplied in a `bind()` call should allow reuse of local addresses.


`SO_KEEPALIVE` enables the periodic transmission of messages on a connected socket. Should the connected party fail to respond to these messages, the connection is considered broken and processes using the socket are notified via a `SIGPIPE` signal.

`SO_DONTROUTE` indicates that outgoing messages should bypass the standard routing facilities. Instead, messages are directed to the appropriate network interface according to the network portion of the destination address.

`SO_LINGER` controls the action taken when unsent messages are queued on socket and a `close()` is performed. If the socket promises reliable delivery of data and `SO_LINGER` is set, the system will block the process on the close attempt until it is able to transmit the data or until it decides it is unable to deliver the information (a timeout period, termed the linger interval, is specified in the `setsockopt` call when `SO_LINGER` is requested). If `SO_LINGER` is disabled and a close is issued, the system will process the close in a manner that allows the process to continue as quickly as possible.

`SO_BROADCAST` requests permission to send broadcast datagrams on the socket. Broadcast packets are only used to send data to groups of recipients.

`SO_OOBINLINE` requests that out-of-band data (with protocols that support out-of-band data) be placed in the normal data input queue as received; it will then be accessible with `recv()` or `read` calls without the `MSG_OOB` flag.



`SO_SNDBUF` and `SO_RCVBUF` are options to adjust the normal buffer sizes allocated for output and input buffers, respectively. The buffer size may be increased for high-volume connections, or may be decreased to limit the possible backlog of incoming data. The system places an absolute limit on these values.

Finally, `SO_TYPE` and `SO_ERROR` are options used only with `getsockopt`.

`SO_TYPE` returns the type of the socket, such as `SOCK_STREAM`; it is useful for servers that inherit sockets on startup.

`SO_ERROR` returns any pending error on the socket and clears the error status. It may be used to check for asynchronous errors on connected datagram sockets or for other asynchronous errors.

RETURN VALUE

A 0 is returned if the call succeeds, if it fails a -1 is returned and a corresponding error code is written in the global variable `errno`.


ERRORS

The call fails if:

[EBADF]	The argument <code>s</code> is not a valid descriptor.
[ENOPROTOOPT]	The option is unknown at the level indicated.

SEE ALSO

`bind()`, `close()`, `getprotobyname()`, `getprotobynumber()`, `getprotoent()`, `ioctl()`, `recv()`, `setsockopt()`, `socket()`



2.10 ioctl()

NAME

ioctl() - set or read device characteristics

SYNOPSIS

#include <ioctl.h>

```
int ioctl(d, request, argp)  
    int d;  
    unsigned long request;  
    char *argp;
```

DESCRIPTION

Ioctl performs a variety of functions on open descriptors. In particular, many operating characteristics of character special files (e.g. terminals) may be controlled with ioctl requests. The write-up of various devices in the manual pages discuss how ioctl applies to them.

An ioctl request has encoded in it whether the argument is an in parameter or out parameter, and the size of the argument argp in bytes. Macros and defines used in specifying an ioctl request are located in the file <ioctl.h>.


RETURN VALUE

A 0 is returned if the call succeeds, if it fails a -1 is returned and a corresponding error code is written in the global variable errno.

ERRORS

The call fails if:

[EBADF]	D is not a valid descriptor.
----------------	------------------------------



2.11 listen()

NAME

listen() - listen for connections on a socket

SYNOPSIS

```
int listen(s, backlog)  
int s, backlog;
```

DESCRIPTION

To accept connections, a socket is first created with `socket()`, a willingness to accept incoming connections and a queue limit for incoming connections are specified with `listen()`, and then the connections are accepted with `accept()`. The `listen` call applies only to sockets of type `SOCK_STREAM` or `SOCK_SEQPACKET`.

The `backlog` parameter defines the maximum length the queue of pending connections may grow to. If a connection request arrives with the queue full, the client may receive an error with an indication of `ECONNREFUSED`, or, if the underlying protocol supports retransmission, the request may be ignored so that retries may succeed.

RETURN VALUE

A 0 is returned if the call succeeds, if it fails a -1 is returned and a corresponding error code is written in the global variable `errno`.

ERRORS

The call fails if:


[EBADF]	The argument <code>s</code> is not a valid descriptor.
[EOPNOTSUPP]	The socket is not of a type that supports the operation <code>listen</code> .

SEE ALSO

`accept()`, `connect()`, `socket()`

NOTE

The `backlog` is currently limited (silently) to 5.



2.12 read()

NAME

read() - read input from a socket

SYNOPSIS

```
int read(socket, buf, len)  
    int socket;  
    char *buf;  
    int len;
```

DESCRIPTION

Read is an equivalent for the normal OS-9 read() - function, and is used in exactly the same manner. It attempts to read len bytes of data from the object referenced by the descriptor socket into the buffer pointed to by buf.

Normally, if no data is ready to be read from the socket, read() waits until data can be read. However, if the socket is marked as non-blocking (see ioctl()), read will not wait - instead it will return a -1 and write the value EWOULDBLOCK in the global variable errno.

The functions _gs_rdy() or select() can be used to determine whether or not data is available to be read at the socket.

RETURN VALUE

If successful, the number of bytes actually read is returned. Otherwise, a -1 is returned and the global variable errno is set to indicate the error.


ERRORS

The call fails if:

[EBADF]	D is not a valid file or socket descriptor open for reading.
[EWOULDBLOCK]	The file was marked for non-blocking I/O, and no data were ready to be read.

SEE ALSO

_gs_rdy(), accept(), ioctl(), recv(), select(), socket(), write()



2.13 recv()

NAME

recv() - receive a message from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int recv(socket, buf, len, flags)
```

```
    int socket;
```

```
    char *buf;
```

```
    int len, flags;
```

DESCRIPTION

recv() is used to receive messages from a socket. The recv() call is normally used only on a connected socket (see connect()).

Socket defines the socket from which the data should be read. The parameter buf is the address of the buffer into which the received data will be written and len is the length of this buffer.

If no messages are available at the socket, the receive call waits for a message to arrive, unless the socket is non-blocking (see ioctl()) in which case a return code of -1 is returned with the global variable errno set to EWOULDBLOCK.


The _gs_rdy() and select() calls may be used to determine when more data arrives.

The flags argument to a recv call is formed by OR'ing one or more of the following values,

```
#define MSG_OOB          0x1    /* process out-of-band data */
#define MSG_PEEK         0x2    /* peek at incoming message */
```

RETURN VALUE

recv() returns the number of bytes successfully received, or, if an error occurred, a return code of -1 is returned with the global variable errno set to the appropriate error code.




ERRORS

The call fails if:

[EBADF]	The argument <i>s</i> is an invalid descriptor.
[EWOULDBLOCK]	The socket is marked non-blocking and the receive operation would block.

SEE ALSO

`_gs_rdy()`, `getsockopt()`, `ioctl()`, `read()`, `recvfrom()`, `recvmsg()`, `select()`, `send()`, `sendmsg()`, `sendto()`, `socket()`



2.14 recvfrom()

NAME

recvfrom() - read data from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int recvfrom(socket, buf, len, flags, from, fromlen)
```

```
    int socket;
```

```
    char *buf;
```

```
    int len, flags;
```

```
    struct sockaddr *from;
```

```
    int *fromlen;
```

DESCRIPTION

recvfrom() is used to receive messages from a socket. Recvfrom may be used to receive data from a socket whether it is in a connected state or not (see connect()).

If a message is too long to fit in the supplied buffer, excess bytes may be discarded depending on the type of socket the message is received from (see socket()).

Socket defines the socket from which the data should be read. The parameter buf is the address of the buffer into which the received data will be written and len is the length of this buffer. If from is non-zero, the source address of the message is filled in. Fromlen is a value-result parameter, initialized to the size of the buffer associated with from, and modified on return to indicate the actual size of the address stored there.

If no messages are available at the socket, the receive call waits for a message to arrive, unless the socket is non-blocking (see ioctl()) in which case a return code of -1 is returned with the global variable errno set to EWOULDBLOCK.

The _gs_rdy() and select() calls may be used to determine when more data arrives.


The flags argument to a recv call is formed by OR'ing one or more of the following values:

```
#define MSG_OOB
```

```
0x1    /* process out-of-band data */
```

```
#define MSG_PEEK
```

```
0x2    /* peek at incoming message */
```



RETURN VALUE

recvfrom() returns the number of bytes successfully received, or, if an error occurred, a return code of -1 is returned with the global variable errno set to the appropriate error code.


ERRORS

The call fails if:

[EBADF]	The argument s is an invalid descriptor.
[EWOULDBLOCK]	The socket is marked non-blocking and the receive operation would block.

SEE ALSO

_gs_rdy(), getsockopt(), ioctl(), read(), recvfrom(), recvmsg(), select(), send(), sendmsg(), sendto(), socket()



2.15 recvmsg()

NAME

recvmsg() - read data from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int recvmsg(socket, msg, flags)
    int socket;
    struct msghdr *msg;
    int flags;
```


DESCRIPTION

recvmsg() is used to receive messages from a socket. recvmsg() may be used to receive data from a socket whether it is in a connected state or not (see connect()).

Socket is the socket from which the data should be read. The parameter msg is a pointer to a msghdr structure that the recvmsg call uses to minimize the number of directly supplied parameters. This structure has the following form, as defined in <socket.h>:

```
struct msghdr {
    caddr_t    msg_name;           /* optional address */
    int        msg_namelen;        /* size of address */
    struct iovec *msg_iov;          /* scatter/gather array */
    int        msg_iovlen;         /* # elements in msg_iov */
    caddr_t    msg_accrights;       /* access rights sent/received */
    int        msg_accrightslen;
};
```

Here msg_name and msg_namelen specify the destination address if the socket is unconnected; msg_name may be given as a null pointer if no names are desired or required. A buffer to receive any access rights sent along with the message is specified in msg_accrights, which has the length msg_accrightslen. The parameter msg_iov is a pointer to the array of iovec structures with msg_iovlen defining the number of structures in the array. The iovec structure is used to hold the addresses to which the received files should be copied. The iovec structure is defined in the uio.h file and has the following form:



```
struct iovec {
    caddr_t      io_base;
    int          iov_len;
};
```

Where:

iov_base is the start address for a memory area which may be used to write the received data.

iov_len is the available length of this memory area in bytes

Each memory area defined in the array is completely filled before the pointer `msg_iov` steps to the next area. Each memory area is used in the order of its appearance in the array.

If no messages are available at the socket, `recvmsg()` waits for a message to arrive, unless the socket is nonblocking (see `ioctl()`) in which case a return code of -1 is returned with the external variable `errno` set to `EWOULDBLOCK`.

The `_gs_rdy()` or `select()` calls may be used to determine when more data has arrived at the socket.

The flags argument to a `recvmsg()` call is formed by OR'ing one or more of the values,

```
#define MSG_OOB          0x1  /* process out-of-band data */
#define MSG_PEEK         0x2  /* peek at incoming message */
```

RETURN VALUE

`recvmsg()` returns the number of bytes successfully received, or, if an error occurred, a return code of -1 is returned with the global variable `errno` set to the appropriate error code.

ERRORS


The calls fail if:

[EBADF] The argument `s` is an invalid descriptor.

[EWOULDBLOCK] The socket is marked non-blocking and the receive operation would block.

SEE ALSO

`getsockopt()`, `ioctl()`, `read()`, `recvfrom()`, `recvmsg()`, `select()`, `send()`, `sendmsg()`, `sendto()`, `setsockopt()`, `socket()`



2.16 select

NAME

select() - synchronous I/O path controlling

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <time_stuff.h>
```

```
int select(nfds, readfds, writefds, exceptfds, timeout)
    int nfds;
    fd_set *readfds, *writefds, *exceptfds;
    struct timeval *timeout;
```

```
FD_SET(fd, &fdset)
FD_CLR(fd, &fdset)
FD_ISSET(fd, &fdset)
FD_ZERO(&fdset)
    int fd;
    fd_set fdset;
```


DESCRIPTION

Select examines the I/O descriptor sets whose addresses are passed in readfds, writefds, and exceptfds to see if some of their descriptors are ready for reading, are ready for writing, or have an exceptional condition pending, respectively. The first nfds descriptors are checked in each set; i.e. the descriptors from 0 through nfds-1 in the descriptor sets are examined. On return, select replaces the given descriptor sets with subsets consisting of those descriptors that are ready for the requested operation. The total number of ready descriptors in all the sets is returned in nfound.

The descriptor sets are stored as bit fields in arrays of integers. The following macros (defined in natdefs.h) are provided for manipulating such descriptor sets:

FD_ZERO(&fdset)	zeros the descriptor set fdset
FD_SET(fd, &fdset)	sets the bit for the descriptor fd in fdset.
FD_CLR(fd, &fdset)	deletes the bit for the descriptor fd from fdset.
FD_ISSET(fd, &fdset)	is nonzero if fd is a member of fdset, zero otherwise.

The behavior of these macros is undefined if a descriptor value is less than zero or greater than or equal to FD_SETSIZE, which is normally at least equal to the maximum number of descriptors supported by the system.



If timeout is a non-zero pointer, it specifies a maximum interval to wait for the selection to complete. If timeout is a zero pointer, the select blocks indefinitely. To affect a poll, the timeout argument should be non-zero, pointing to a zero-valued timeval structure.

Any of readfds, writefds, and exceptfds may be given as zero pointers if no descriptors are of interest.

RETURN VALUE

Select returns the number of ready descriptors that are contained in the descriptor sets, or -1 if an error occurred. If the time limit expires then select returns 0. If select returns with an error, including one due to an interrupted call, the descriptor sets will be unmodified.

ERRORS

The call fails if:


[EBADF]	One of the descriptor sets specified an invalid descriptor.
[EINTR]	A signal was delivered before the time limit expired and before any of the selected events occurred.
[EINVAL]	The specified time limit is invalid. One of its components is negative or too large.

NOTE

It is possible to select I/O paths which are not sockets. This is limited to those paths where `_gs_rdy()` may also be used. Such non-socket I/O paths may only be selected for reading.

SEE ALSO

`accept()`, `connect()`, `_gs_rdy()`, `ioctl()`, `read()`, `recv()`, `recvfrom()`, `recvmsg()`, `select()`, `send()`, `sendmsg()`, `sendto()`, `socket()`, `write()`



2.17 send()

NAME

send() - send a message from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int send(socket, buf, len, flags)
```

```
    int socket;
```

```
    char *buf;
```

```
    int len, flags;
```

DESCRIPTION

send() is used to transmit a message to another socket. send() may be used only when the socket is in a connected state.

Socket defines the socket from which the data should be read. The parameter buf is the address of the buffer from which the to be transmitted data will be read and len is the length of this buffer.

If no message space is available at the socket to hold the message to be transmitted, send() will normally block. However, if the socket has been placed in non-blocking I/O mode, send() will return a -1 and write the value EWOULDBLOCK in the global variable errno. The select() call may be used to determine when it is possible to send more data.


The flags parameter may include one or more of the following:

```
#define MSG_OOB          0x1    /* process out-of-band data */
```

```
#define MSG_DONTROUTE    0x4    /* bypass routing, use direct interface */
```

MSG_OOB is used to send out-of-band data on sockets that support this notion (e.g. SOCK_STREAM); the underlying protocol must also support out-of-band data.

MSG_DONTROUTE is usually used only by diagnostic or routing programs.



RETURN VALUE

The call returns the number of characters successfully sent, or, if an error occurred it will return a -1 and write the appropriate error code in the global variable `errno`.


ERRORS

The call fails if:

[EBADF]	An invalid descriptor was specified.
[EMSGSIZE]	The socket requires that message be sent atomically, and the size of the message to be sent made this impossible.
[EWOULDBLOCK]	The socket is marked non-blocking and the requested operation would block.

SEE ALSO

`getsockopt()`, `ioctl()`, `recv()`, `recvfrom()`, `recvmsg()`, `select()`, `sendmsg()`, `sendto()`, `setsockopt()`, `socket()`, `write()`



2.18 sendx()

NAME

sendx() – send a message from a socket (available only on ETH29/FDDI29)

SYNOPSIS

```
int sendx(socket, sbuf, lbuf, length, flags)  
int socket;  
u_char *sbuf, *lbuf;  
in length, flags;
```

DESCRIPTION


The sendx() call works in the same manner as the send() call, with one important difference: the address of the sendbuf is local to the FDDI29/ETH29's memory and not local to the host's memory in case of the send() call. Buffer space is allocated using the SysReqMem() call, it is returned with the SysRetMem() call.

Flags could be:

MSG_OOB	0x01
MSG_DONTROUTE	0x04
MSG_SWAPWORD	0x8000
MSG_SWAPBYTE	0x4000
MSG_USEDMA	0x2000

SEE ALSO

send(), SysReqMem(), SysRetMem()



2.19 sendmsg()

NAME

sendmsg() - send a message from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int sendmsg(socket, msg, flags)
    int socket;
    struct msghdr *msg[];
    int flags;
```


DESCRIPTION

sendmsg() is used to transmit a message to another socket. Sendmsg() may be used at any time regardless of whether the socket is connected or not (see connect()).

Socket defines the socket from which the data should be read. The parameter msg is a pointer to the msghdr structure into which the received data will be written. The parameter msg is used to reduce the number parameters to be passed. The structure msghdr is defined in the socket.h file and has the following form:

```
struct msghdr {
    caddr_t    msg_name;           /* optional address */
    int        msg_namelen;        /* length of the optional address */
    struct iovec *msg_iov;          /* scatter/gather array */
    int        msg_iovlen;         /* # elements in msg_iov */
    caddr_t    msg_accrights;      /* access rights */
    int        msg_accrightslen;
};
```

Here msg_name and msg_namelen specify the destination address if the socket is unconnected; msg_name may be given as a null pointer if no names are desired or required. A buffer to receive any access rights sent along with the message is specified in msg_accrights, which has the length msg_accrightslen. The parameter msg_iov is a pointer to the array of iovec structures with msg_iovlen defining the number of structures in the array. The iovec structures is used hold the addresses from which the to be transmitted files should be copied. The iovec structure is defined in the uio.h file and has the following form:



```
struct iovec {
    caddr_t    io_base;
    int        iov_len;
};
```

Where:

iov_base is the start address for a memory area which may be used to read the to be transmitted data.

iov_len is the available length of this memory area in bytes

Each memory area defined in the array is completely read before the pointer `msg_iov` steps to the next area. Each memory area is read in the order of its appearance in the array.

If not enough message space is available at the socket to hold the message to be transmitted, then `sendmsg()` normally blocks, unless the socket has been placed in non-blocking I/O mode. The `select()` call may be used to determine when it is possible to send more data.

The flags parameter may include one or more of the following:


```
#define MSG_OOB          0x1    /* process out-of-band data */
#define MSG_DONTROUTE    0x4    /* bypass routing, use direct interface */
```

MSG_OOB is used to send out-of-band data on sockets that support this notion (e.g. `SOCK_STREAM`); the underlying protocol must also support out-of-band data.

MSG_DONTROUTE is usually used only by diagnostic or routing programs.

RETURN VALUE

The call returns the number of characters successfully sent, or, if an error occurred it will return a -1 and write the appropriate error code in the global variable `errno`.




ERRORS

The call fails if:

[EBADF]	An invalid descriptor was specified.
[EMSGSIZE]	The socket requires that message be sent atomically, and the size of the message to be sent made this impossible.
[EWOULDBLOCK]	The socket is marked non-blocking and the requested operation would block.

SEE ALSO

getsockopt(), ioctl(), recv(), recvfrom(), recvmsg(), select(), send(), sendto(), setsockopt(), socket(), write()



2.20 sendto()

NAME

sendto() - send a message from a socket

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int sendto(socket, buf, len, flags, to, tolen)
```

```
    int socket;
```

```
    char *buf;
```

```
    int len, flags;
```

```
    struct sockaddr *to;
```

```
    int tolen;
```

DESCRIPTION

sendto() is used to transmit a message to another socket. sendto() may be used at any time regardless of whether a socket is connected or not (see connect()).

Socket defines the socket from which the data should be read. The parameter buf is the address of the buffer from which the data will be transmitted and len is the length of this buffer. The socket address of the target is given by to with tolen specifying its size. If the message is too long to pass atomically through the underlying protocol, then the error EMSGSIZE is returned, and the message is not transmitted.

If no message space is available at the socket to hold the message to be transmitted, then sendto() normally blocks, unless the socket has been placed in non-blocking I/O mode. The select() call may be used to determine when it is possible to send more data.


The flags parameter may include one or more of the following:

```
#define MSG_OOB          0x1    /* process out-of-band data */
```

```
#define MSG_DONTROUTE    0x4    /* bypass routing, use direct interface */
```

MSG_OOB is used to send out-of-band data on sockets that support this notion (e.g. SOCK_STREAM); the underlying protocol must also support out-of-band data.

MSG_DONTROUTE is usually used only by diagnostic or routing programs.



RETURN VALUE

The call returns the number of characters successfully sent, or, if an error occurred it will return a -1 and write the appropriate error code in the global variable `errno`.


ERRORS

The call fails if:

[EBADF]	An invalid descriptor was specified.
[EMSGSIZE]	The socket requires that message be sent atomically, and the size of the message to be sent made this impossible.
[EWOULDBLOCK]	The socket is marked non-blocking and the requested operation would block.

SEE ALSO

`getsockopt()`, `ioctl()`, `recv()`, `recvfrom()`, `recvmsg()`, `select()`, `send()`, `sendmsg()`, `setsockopt()`, `socket()`, `write()`



2.21 shutdown()

NAME

shutdown() - shut down part of a full-duplex connection

SYNOPSIS

```
int shutdown(socket, how)  
int socket, how;
```

DESCRIPTION

shutdown() causes all or part of a full-duplex connection on the socket associated with socket to be shut down.

If how is 0, then further receives will be disallowed. If how is 1, then further sends will be disallowed. If how is 2, then further sends and receives will be disallowed. A send or receive function that attempts to access the socket after it has been shutdown will return an error code.

RETURN VALUE

A 0 is returned if the call succeeds, -1 if it fails and a corresponding error code is written to the global variable errno.


ERRORS

The call fails if:

[ENOTCONN]	The specified socket is not connected.
------------	--

SEE ALSO

connect(), socket()



2.22 setsockopt()

NAME

setsockopt() - set socket options

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
setsockopt(socket, level, optname, optval, optlen)
           int socket, level, optname;
           char *optval;
           int optlen;
```


DESCRIPTION

setsockopt() is used to set the options associated with a socket. Options may exist at multiple protocol levels; they are always present at the uppermost ``socket" level.

When setting socket options, the level at which the option resides and the name of the option must be specified. To set options at the ``socket" level, level is specified as SOL_SOCKET. To manipulate options at any other level the protocol number of the appropriate protocol controlling the option must be entered. For example, to indicate that an option is to be interpreted by the TCP protocol, level should be set to the protocol number of TCP; see getprotoent().

The parameters optval and optlen are used by setsockopt to pass option values. If no option value is to be set, optval may be given as 0.

Optname and any specified options are passed uninterpreted to the appropriate protocol module for interpretation. The include file <socket.h> contains definitions for ``socket" level options, described below. Options at other protocol levels vary in format and name; consult the corresponding entries in the appropriate manuals.




2.22.1 Socket level options


Most socket-level options take an `int` parameter for `optval`. For `setsockopt`, the parameter should non-zero to enable a boolean option, or zero if the option is to be disabled. `SO_LINGER` uses a pointer to a `linger` structure parameter that is defined in `<socket.h>`, and which specifies the desired state of the option and the linger interval (see below).

The following options are recognized at the socket level. Except as noted, each may be examined with `getsockopt` and set with `setsockopt`.

SO_DEBUG	toggle recording of debugging information
SO_REUSEADDR	toggle local address reuse
SO_KEEPALIVE	toggle keep connections alive
SO_DONTROUTE	toggle routing bypass for outgoing messages
SO_LINGER	linger on close if data present
SO_BROADCAST	toggle permission to transmit broadcast messages
SO_OOBINLINE	toggle reception of out-of-band data in band
SO_SNDBUF	set buffer size for output
SO_RCVBUF	set buffer size for input
SO_DEBUG	enables debugging in the underlying protocol modules. The data is written to specific address areas and can later be viewed using the program <code>netprint</code> .
SO_REUSEADDR	indicates that the rules used in validating addresses supplied in a <code>bind()</code> call should allow reuse of local addresses
SO_KEEPALIVE	enables the periodic transmission of messages on a connected socket. Should the connected party fail to respond to these messages, the connection is considered broken and processes using the socket are notified via a <code>SIGPIPE</code> signal.
SO_DONTROUTE	indicates that outgoing messages should bypass the standard routing facilities. Instead, messages are directed to the appropriate network interface according to the network portion of the destination address.



SO_LINGER	controls the action taken when unsent messages are queued on socket and a close() is performed. If the socket promises reliable delivery of data and SO_LINGER is set, the system will block the process on the close attempt until it is able to transmit the data or until it decides it is unable to deliver the information (a timeout period, termed the linger interval, is specified in the setsockopt call when SO_LINGER is requested). If SO_LINGER is disabled and a close is issued, the system will process the close in a manner that allows the process to continue as quickly as possible.
SO_BROADCAST	requests permission to send broadcast datagrams on the socket. Broadcast packets are only used to send data to groups of recipients.
SO_OOINLINE	requests that out-of-band data (with protocols that support out-of-band data) be placed in the normal data input queue as received; it will then be accessible with recv() or read calls without the MSG_OOB flag.
SO_SNDBUF SO_RCVBUF	and are options to adjust the normal buffer sizes allocated for output and input buffers, respectively. The buffer size may be increased for high-volume connections, or may be decreased to limit the possible backlog of incoming data. The system places an absolute limit on these values.




2.22.2 IP multicasting

The IP multicasting capability has been implemented according to the RFC 1112. Currently it is available only for the ETH29 and FDDI29 boards. Of course, IP multicasting is not possible on ISO/OSI.

To enable IP multicasting, the parameter level has to be set to the protocol number of IP which is IPPROTO_IP.

There are five options which concern the IP multicasting feature; they are passed in the optname parameter. Here is a detailed list of these options

IP_MULTICAST_IF	Enables the multicasting capability on a given interface. The IP address of this interface is passed in the optval parameter which is a pointer to a struct in_addr structure (defined in the <netinet/in.h> include file). IP multicasting is disabled when the node has left the last multicasting group.
IP_ADD_MEMBERSHIP	Adds a membership of this host to a specified multicasting group. The own IP address and the multicasting address of the group to be joined are specified in a struct ip_mreq structure, a pointer to which is passed in the optval parameter. The struct ip_mreq structure is defined in the <netinet/in.h> include file.
IP_DROP_MEMBERSHIP	Deletes the membership of this host from a specified multicasting group. As in the IP_ADD_MEMBERSHIP call, the group is specified in a struct ip_mreq structure, a pointer to which is passed in the optval parameter.
IP_MULTICAST_TTL	Used to change the Time To Live value of IP multicasting packets. The Time To Live value specifies how many gateways a multicasting message may pass and thus how far it will get in the network topography. The default Time To Live value is 1, so it has to be explicitly set to a higher value if multicast messages are supposed to reach hosts beyond the own network. The optval parameter is a pointer to an u_char containing the desired Time To Live value.



IP_MULTICAST_LOOP

This option is used to specify whether or not IP multicasting messages are sent to the loopback interface. By default they are sent to the loopback interface since the sending host is itself a member of the corr. multicast group. The optval parameter is a pointer to an u_char containing TRUE (send to loopback interface) or FALSE (do not send to loopback interface).

RETURN VALUE

A 0 is returned if the call succeeds, -1 if it fails and a corresponding error code is written to the global variable errno.

ERRORS

The call fails if:

[EBADF]


The argument s is not a valid descriptor.

[ENOPROTOOPT]

The option is unknown at the level given.

SEE ALSO

bind(), close(), getprotobyname(), getprotobyname(), getprotoent(3N), getsockopt(), ioctl(), recv(), socket(),



2.23 socket()

NAME

socket() - create an endpoint for communication

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <socket.h>
```

```
int socket(domain, type, protocol)
           int domain, type, protocol;
```

DESCRIPTION

socket() creates an endpoint for communication and returns an integer as a descriptor for it.

The domain parameter specifies a communications domain within which communication will take place; this selects the protocol family which should be used. The protocol family generally is the same as the address family for the addresses supplied in later operations on the socket. These families are defined in the include file <socket.h>. The currently understood format is


AF_INET (ARPA Internet protocols)

AF_ISO (ISO/OSI)

The socket has the indicated type, which specifies the semantics of communication. The currently defined types are:

SOCK_STREAM	type provides sequenced, reliable, two-way connection based byte streams. An out-of-band data transmission mechanism may be supported.
SOCK_DGRAM	socket supports datagrams (connectionless, unreliable messages of a fixed (typically small) maximum length).
SOCK_RAW	sockets provide access to internal network protocols and interfaces. The type SOCK_RAW, which is available only to the super-user, is not described here.
SOCK_SEQPACKET	sequenced packet protocol; this must be used instead of SOCK_STREAM when working on ISO/OSI.

The protocol specifies which protocol is to be used with the socket. Normally only a




single protocol exists to support a particular socket type within a given protocol family. However, it is possible that many protocols may exist, in which case a particular protocol must be specified in this manner. The protocol number to use is particular to the communication domain in which communication is to take place; see protocols(3N).

Sockets of type `SOCK_STREAM` are full-duplex byte streams, similar to pipes. A stream socket must be in a connected state before any data may be sent or received on it. A connection to another socket is created with a `connect()` call. Once connected, data may be transferred using `read()` and `write()` calls or some variant of the `send()` and `recv()` calls. When a session has been completed a `close()` may be performed. Out-of-band data may also be transmitted as described in `send()` and received as described in `recv()`.

The communications protocols used to implement a `SOCK_STREAM` insure that data is not lost or duplicated. If a piece of data for which the peer protocol has buffer space cannot be successfully transmitted within a reasonable length of time, then the connection is considered broken and calls will indicate an error with -1 returns and with `ETIMEDOUT` as the specific code in the global variable `errno`.

`SOCK_DGRAM` and `SOCK_RAW` sockets allow sending of datagrams to correspondents named in `sendto()` or `sendmsg()` calls. The `send()` and `write()` functions can be used with these types of sockets, if the destination is first set using `connect()`. Datagrams are generally received with `recvfrom()`, which returns the next datagram along with its return address.

The type `SOCK_STREAM` which is used in the TCP/IP stack for connection oriented sockets must be replaced by the `SOCK_SEQPACKET` type (sequenced packet protocol) when working on ISO/OSI.



The function `ioctl()` supports the command `FIONBIO` which can be used to toggle the socket between blocking and non-blocking mode.

```
#include <ioctl.h>
...
int s;
int nonblock = 1;
...
s = socket(AF_INET, SOCK_STREAM, 0);
...
if (ioctl(s, FIONBIO, (unsigned char *)&nonblock) < 0) {
    perror("ioctl (FIONBIO)");
    exit(1);
}
```

By setting the variable `nonblock` to 0 the socket can be toggled back to the default setting "blocking".

The operation of sockets is controlled by socket level options. These options are defined in the file `<socket.h>`. `setsockopt()` and `getsockopt()` are used to set and get options, respectively.

RETURN VALUE

A -1 is returned if an error occurs and a corresponding error code is written to the global variable `errno`. Otherwise, the return value is a descriptor referencing the socket.


ERRORS

The call fails if:

[EPROTONOSUPPORT]	The protocol type or the specified protocol is not supported within this domain.
[EACCESS]	Permission to create a socket of the specified type and/or protocol is denied.
[ENOBUFFS]	Insufficient buffer space is available. The socket cannot be created until sufficient resources are freed.

SEE ALSO

`accept()`, `bind()`, `connect()`, `getsockname()`, `getsockopt()`, `ioctl()`, `listen()`, `read()`, `recv()`, `recvfrom()`, `recvmsg()`, `select()`, `send()`, `sendmsg()`, `sendto()`, `setsockopt()`, `shutdown()`, `write()`



2.24 SysReqMem()

NAME

SysReqMem() – allocate system memory on the FDDI29/ETH29 slave board

SYNOPSIS

```
char *SysReqMem(socket, memsize)  
int socket;  
u_long memsize;
```

DESCRIPTION


The SysReqMem() call is used to allocate system memory on the FDDI29/ETH29 slave board. The socket number supplied to the call is the one returned from a previous socket() call. The parameter memsize gives the requested memory size.

RETURN VALUE

Two return values are possible: 0, if no memory could be allocated in the desired size, or the address of the allocated memory.

SEE ALSO

send(), sendx(), SysRetMem()



2.25 SysRetMem()

NAME

SysRetMem() – return system memory previously allocated by SysReqMem() on the FDDI29/ETH29 slave board

SYNOPSIS

```
int SysRetMem(socket, memsize, memadr)  
int socket;  
u_long memsize;  
char *memadr;
```


DESCRIPTION

The SysRetMem() call is used to return the system memory allocated previously by a SysReqMem() call. The parameters supplied to the call are the socket number and the memory size as already used by the SysReqMem() call, and the buffer address as returned by the SysReqMem() call.

SysRetMem() has to be called before a socket is closed. This is due to the fact that memory previously allocated by SysReqMem() is by no means related to any socket and therefore is not automatically freed when the socket is closed. But as the socket number is the only way to communicate with the slave board, you cannot free memory after you have closed the socket.

SEE ALSO

send(), sendx(), SysReqMem()



2.26 write()

NAME

write() - write output

SYNOPSIS

```
int write(socket, buf, len)  
    int socket;  
    char *buf;  
    int len;
```

DESCRIPTION

write() is an exact equivalent for the normal OS-9 write() - function and functions precisely the same. It attempts to write len bytes of data to the object referenced by the descriptor socket from the buffer pointed to by buf.

The parameter socket defines the socket from which the data should be read. The socket was created using either socket() or accept(). The parameter buf is the address of the buffer from which the to be transmitted data will be read and len is the length of this buffer.

If no message space is available at the socket to hold the message to be transmitted, write() will normally block. However, if the socket has been placed in non-blocking I/O mode, write() will return a -1 and write the value EWOULDBLOCK in the global variable errno. The select() call may be used to determine when it is possible to send more data.

If write() fails, the file pointer will remain unchanged.

RETURN VALUE

Upon successful completion, the number of bytes actually written is returned. Otherwise a -1 is returned and the global variable errno is set to indicate the error.


ERRORS

The call fails if:

[EBADF]	Socket is not a valid descriptor open for writing.
[EWOULDBLOCK]	The Socket was marked for non-blocking I/O, and no data could be written immediately.

SEE ALSO


accept(), ioctl(), read(), select(), send(), socket()



3 Conversion Routines

The following pages contain a description of the byte-order conversion utilities included in this package. These routines are provided as macros in the file `in.h` and can be used to convert words between the host and network ordering of bytes.

htonl()	convert long word from host to network
htons()	convert short word from host to network
ntohl()	convert long word from network to host
ntohs()	convert short word from network to host



3.1 htonl()

NAME

htonl() - Convert Long Word : Host -> Network

DEFINITION

#include <in.h>

unsigned long htonl (hostlong)
unsigned long hostlong

DESCRIPTION

The function htonl() converts a long word from the ordering used by the system where htonl() was called to the byte ordering required by the network.


For systems with 680x0 processors, this function is available as a macro in the file in.h.

RESULTS

The word passed is returned in the network byte ordering.

SEE ALSO

gethostbyname(), getservent(), htons(), ntohl(), ntohs()



3.2 htons()

NAME

htons() - Convert Short Word : Host -> Network

DEFINITION

#include <in.h>

unsigned short htons (hostshort)
unsigned short hostshort

DESCRIPTION

The function htons() converts a short word from the ordering used by the system where htons() was called to the byte ordering required by the network.


For systems with 680x0 processors, this function is available as a macro in the file in.h.

RESULTS

The word passed is returned in the network byte ordering.

SEE ALSO

gethostbyname(), getservent(), htons(), ntohl(), ntohs()



3.3 ntohs()

NAME

ntohl() - Convert Long Word : Network -> Host

DEFINITION

#include <in.h>

unsigned long ntohs (netlong)
unsigned long netlong

DESCRIPTION

The function ntohs() converts a long word from the ordering used by the network to the byte ordering required by the system where ntohs() was called.


For systems with 680x0 processors, this function is available as a macro in the file in.h.

RESULTS

The word passed is returned in the byte ordering used internal to the system.

SEE ALSO

gethostbyname(), getservent(), htons(), ntohs(), ntohs()



3.4 ntohs()

NAME

ntohs() - Convert Short Word : Network -> Host

DEFINITION

#include <in.h>

unsigned long ntohs (netshort)
unsigned long netshort

DESCRIPTION

The function ntohs() converts a short word from the ordering used by the network to the byte ordering required by the system where ntohl() was called.


For systems with 680x0 processors, this function is available as a macro in the file in.h.

RESULTS

The word passed is returned in the byte ordering used internal to the system.

SEE ALSO

gethostbyname(), getservent(), htons(), ntohl(), ntohs()




4 Address Manipulation Functions

The functions described in this chapter enable the manipulation of the Internet addresses. The functions can be found in the library `inetdb.1`.

Currently the ISO/OSI versions of these calls are not yet implemented. But in the near future, the address manipulating functions will be available for both TCP/IP and ISO/OSI, on ISO/OSI they begin with a preceding `iso_` instead of `inet_` than. It is necessary to differentiate in this way since the address formats used by both protocols are different, too. The functionality, however, is the same on both protocols.

These are the TCP/IP versions of the calls:

<code>inet_addr()</code>	Generates a complete numerical Internet address from a string containing an Internet address in standard dot notation.
<code>inet_network()</code>	Generates a complete numerical Internet network number from a string containing an Internet address in standard dot notation.
<code>inet_ntoa()</code>	Generates a string in standard dot notation from a numerical Internet address.
<code>inet_makeaddr()</code>	Generates a complete numerical Internet address by combining a network number with a numerical host address.
<code>inet_lnaof()</code>	Extract the numerical host address from a numerical Internet address.
<code>inet_netof()</code>	Extract the network number from a numerical Internet address.




4.1 Internet Addresses in Dot Notation

Internet addresses represented in dot notation can take one of the following four forms:

a.b.c.d	If four parts are given, each part will be interpreted as a single byte. Each part (byte) will then be ordered from left to right to the four bytes of an Internet address.
a.b.c	If three parts are given, the right-most part (c) will be interpreted as a 16-bit unit and will be assigned to the two right-most bytes of the internet address. This agrees with the class B network address form 128.Net.Host.
a.b	If two parts are given, the last part (b) will be interpreted as a 24-bit unit and will be assigned to the 3 right-most bytes of the Internet address. This agrees with the class A network address form Net.Host.
a	If only a single part is given it will be treated as a direct entry of the Internet address.

In dot notation all the parts of an Internet address can be entered in decimal, octal or hexadecimal format. The syntax is the same as for the "C" programming language, thus:

- a leading 0x or 0X signifies that the next digit should be interpreted as **hexadecimal**,
- a leading 0 (zero) signifies that the next digit should be interpreted as **octal**, and
- all other cases should be treated as being in **decimal** notation.



4.2 inet_addr()

NAME

inet_addr() - Internet Address Dot -> Numerical

DEFINITION

```
#include <socket.h>
```

```
#include <in.h>
```

```
#include <inet.h>
```

```
unsigned long inet_addr(cp)
```

```
char *cp;
```

DESCRIPTION


The function `inet_addr()` generates a complete numerical Internet address from a string containing an Internet address in standard dot notation. The parameter `cp` is a pointer to the Internet address string in dot notation. The bytes of the returned numerical Internet address are ordered as required by the network (from left to right).

RESULTS

Successful execution will return the numerical Internet address. If the Internet address in dot-notation was in the incorrect form, a -1 will be returned.

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_network()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_lnaof()`, `inet_netof()`



4.3 inet_lnaof()

NAME

inet_lnaof() - Extract Host Address from Numerical Internet Address

DEFINITION

```
#include <socket.h>
```

```
#include <in.h>
```

```
#include <inet.h>
```

```
unsigned long inet_lnaof(in)
```

```
struct in_addr in;
```

DESCRIPTION


The function `inet_lnaof()` extracts the numerical host address from a numerical Internet address. The function returns the host address in the appropriate byte order for the system which called it.

RESULTS

Successful execution will return the numerical Internet host address. If the Internet address in dot-notation was in the incorrect form, a -1 will be returned.

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_addr()`, `inet_network()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_netof()`



4.4 inet_makeaddr()

NAME

inet_makeaddr - Network + Host Address -> Numerical Internet Address

DEFINITION

```
#include <socket.h>
```

```
#include <in.h>
```

```
#include <inet.h>
```

```
struct in_addr inet_addr(net, lna)
```

```
int net, lna;
```

DESCRIPTION


The function `inet_makeaddr()` generates a complete numerical Internet address by combining a network number with a numerical host address. The bytes of the returned numerical Internet address are ordered as required by the network (from left to right).

RESULTS

Successful execution will return the numerical Internet address in the structure `in_addr` (the 4 bytes are ordered from left to right).

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_network()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_lnaof()`, `inet_netof()`



4.5 inet_netof()

NAME

inet_netof() - Extract Network Address from Numerical Internet Address

DEFINITION

```
#include <socket.h>
#include <in.h>
#include <inet.h>
```

```
unsigned long inet_netof(in)
    struct in_addr in;
```

DESCRIPTION


The function `inet_netof()` extracts the network number from a numerical Internet address. The bytes of the returned network number are ordered as required by the calling system.

RESULTS

Successful execution will return the network number. If the Internet address in dot-notation was in the incorrect form, a -1 will be returned.

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_addr()`, `inet_network()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_lnaof()`



4.6 inet_network()

NAME

inet_network() - Dot Network Address -> Numerical

DEFINITION

```
#include <socket.h>
```

```
#include <in.h>
```

```
#include <inet.h>
```

```
unsigned long inet_network(cp)
```

```
char *cp;
```

DESCRIPTION


The function `inet_network()` generates a complete numerical Internet network number from a string containing an Internet address in standard dot notation. The parameter `cp` is a pointer to the string containing the Internet network address in dot notation. The bytes of the returned numerical Internet address are ordered as required by the calling system.

RESULTS

Successful execution will return the numerical network address. If the Internet network address in dot-notation was in the incorrect form, a -1 will be returned.

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_addr()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_lnaof()`, `inet_netof()`



4.7 inet_ntoa()

NAME

inet_ntoa() - Numerical Internet Address -> Dot

DEFINITION

```
#include <socket.h>
```

```
#include <in.h>
```

```
#include <inet.h>
```

```
char *inet_ntoa(in)  
    struct in_addr in;
```

DESCRIPTION


The function `inet_ntoa()` generates a string in standard dot notation from a numerical Internet address. The parameter `cp` is a pointer to the Internet address string in dot notation. The bytes of the returned numerical Internet address are ordered as required by the network (from left to right).

RESULTS

Successful execution will return the numerical Internet address. If the Internet address in dot-notation was in the incorrect form, a -1 will be returned.

SEE ALSO

`gethostbyname()`, `hosts()`, `inet_addr()`, `inet_network()`, `inet_ntoa()`, `inet_makeaddr()`, `inet_lnaof()`, `inet_netof()`



5 Network Data Basis

The following pages detail the files and functions used for accessing the network data basis. The functions can be found in the library netdb.1

The network data basis consists of the following files:

hosts	data basis for host names
networks	data basis for network numbers
protocols	data basis for protocol names
services	data basis for service names

The functions which can be used to view the data basis are:

To read host entries


gethostbyname()
gethostbyaddr()
gethostent()
sethostent()
endhostent()

To read network entries

getnetbyaddr()
getnetbyname()
getnetent()
setnetent()
endnetent()

To read protocol entries


getprotobyaddr()
getprotobyname()
getprotoent()
setprotoent()
endprotoent()



To read services entries

```
getservbyaddr()  
getservbyname()  
getservbyport()  
getservent()  
setservent()  
endservent()
```

So far, on ISO/OSI only the `gethostbyname()` call has been implemented.



5.1 hosts

NAME

hosts - host name data base

DESCRIPTION

The hosts file contains information regarding the known hosts on the network. For each host a single line should be present with the following information:

- official host name
- Internet address
- aliases

Items are separated by any number of blanks and/or tab characters. A ``#" indicates the beginning of a comment; characters up to the end of the line are not interpreted by routines which search the file.

Network addresses must be entered in the conventional Internet dot notation. Please refer to section 4.1, "Internet Addresses in Dot Notation".

Host names may contain any printable character other than a blank, newline, or comment character. If you wish to be connected into the worldwide Internet network, the hosts file should be based on the official host data base that is administered by the Network Information Control Center (NIC).

FILES


/DD/ETC/hosts

EXAMPLE

```
# The official names are out of the greek, the further names out of the roman mythology
# The names and Internet addresses shown here are not from the official NIC host data
base
127.1          local          localhost lh
192.1.1.1      zeus           jupiter
192.1.1.2      hera           juno
192.1.1.3      hermes          mercury
192.1.1.4      odin            wotan      # in between a germanic
192.1.1.5      ares            mars       # back to greek/roman
192.1.1.6      aphrodite       venus      # the second "token woman"
```

SEE ALSO

ifconfig(), gethostbyaddr(), gethostbyname(), sethostent(), gethostent(), endhostent()



5.1.1 The struct hostent structure

NAME

struct hostent – the host entry structure


DESCRIPTION

The struct hostent structure is used by the functions concerned with the hosts database and describes an internet host. It is defined as follows:

```
struct hostent {
    char        *h_name;           /* official name of host */
    char        **h_aliases;       /* alias list */
    int         h_addrtype;        /* host address type */
    int         h_length;          /* length of address */
    char        **h_addr_list;     /* list of addresses from name server */
};
#define h_addr h_addr_list[0]    /* address, for backward compatibility */
```

The members of this structure are:

h_name	Official name of the host.
h_aliases	A zero terminated array of alternate names for the host.
h_addrtype	The type of address being returned; currentl always AF_INET.
h_length	The length, in bytes, of the address.
h_addr_list	A zero terminated array of network addresses for the host. Host addresses are returned in network byte order.
h_addr	The first address in h_addr_list; this is for backward compatibility.



5.1.2 gethostbyaddr()

NAME

gethostbyaddr() - get network host entry

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <netdb.h>
```

```
extern int h_errno;
```

```
struct hostent *gethostbyaddr(addr, len, type)
    char *addr;
    int len;
    int type;
```

DESCRIPTION

The function gethostbyaddr() takes the address of the desired host in the parameter addr, which has a length of len bytes and is of address type type (currently type can only be AF_INET).

RESULTS

When gethostbyaddr() has been executed successfully, it returns a pointer to a struct hostent structure describing the referenced host (see section 5.1.1, "The struct hostent structure").

Otherwise, an error return status from gethostbyaddr is indicated by return of a null pointer and the external variable h_errno will contain the associated error code.

ERRORS

The call fails if:


HOST_NOT_FOUND The specified host address is unknown.

SEE ALSO

gethostbyname(), sethostent(), gethostent(), endhostent(), hosts

NOTE

All information is contained in a static area so it must be copied, if it is to be saved. Only the Internet address format is currently understood.



5.1.3 gethostbyname()

NAME

gethostbyname() - get network host entry

SYNOPSIS

```
#include <natdefs.h>
#include <netdb.h>
```

```
extern int h_errno;
struct hostent *gethostbyname(name)
    char *name;
```

DESCRIPTION

The function gethostbyname() takes the name of the desired host in the parameter name.

DIFFERENCES ON ISO/OSI

On account of the mass of possible OSI address formats and the problems incident to that there currently is no standard which makes possible a homogeneous administration of addresses and host names, as in the case of TCP/IP with the file /dd/etc/hosts. The N.A.T. GmbH has developed a new gethostbyname() function for ISO/OSI in order to make the address administration as transparent as possible to the user of the N.A.T. OSI/TP protocol software; this function extracts the OSI-address of a given host name out of the /dd/etc/hosts file.


The gethostbyname() function call has a different syntax for the protocol family AF_ISO which is responsible for OSI/TP:

```
struct hostent *gethostbyname(name, addr_fam)
    char *name;
    int *addr_fam;
```

where addr_fam has to be set to AF_ISO.

RESULTS

When gethostbyname() has been executed successfully, it returns a pointer to a struct hostent structure describing the referenced host (see section 5.1.1, "The struct hostent structure"). Otherwise, an error return status from gethostbyname is indicated by return of a null pointer and the external variable h_errno will contain the associated error code.



ERRORS

The call fails if:


HOST_NOT_FOUND The specified host name is unknown.

SEE ALSO

gethostbyaddr(), sethostent(), gethostent(), endhostent(), hosts

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Currently the AF_INET and AF_ISO address format are understood, but the syntax of the call is different for these formats.



5.1.4 gethostent()

NAME

gethostent() - get network host entry

SYNOPSIS

```
#include <natdefs.h>
#include <netdb.h>
```

```
extern int h_errno;
struct hostent *gethostent()
```

DESCRIPTION

The function gethostent() always gets the next entry in the hosts data base. The data base will be opened automatically if necessary.

RESULTS

When gethostent() has been executed successfully, it returns a pointer to a struct hostent structure (see section 5.1.1, "The struct hostent structure"). Otherwise, an error return status from gethostent() is indicated by return of a null pointer and the external variable h_errno will contain the associated error code.

ERRORS

The call fails if:


HOST_NOT_FOUND The specified host name is unknown.

SEE ALSO

gethostbyaddr(), gethostbyname(), sethostent(), endhostent(), hosts

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet address format is currently understood.



5.1.5 sethostent()

NAME

sethostent() - open network hosts data base

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <netdb.h>
```

```
sethostent(stayopen)
```

```
int stayopen;
```


DESCRIPTION

The function sethostent() opens the hosts data base and positions an internal pointer to the first entry in this data base. A subsequent call to gethostent() will return a pointer to the hostent structure for the first entry in this data base (see section 5.1.1, "The struct hostent structure").

If the parameter stayopen contains a non-zero value, the hosts data base will remain open following subsequent calls to gethostbyaddr() or gethostbyname(). Normally, the hosts data base is closed again following every function call.

SEE ALSO

gethostbyaddr(), gethostbyname(), gethostent(), endhostent(), hosts



8.1.5 endhostent()

NAME

endhostent() - Close the hosts data base

SYNOPSIS

```
#include <natdefs.h>
```

```
#include <netdb.h>
```


```
endhostent()
```

DESCRIPTION

The function endhostent() closes the hosts data base after use.

SEE ALSO

gethostbyaddr(), gethostbyname(), sethostent(), gethostent(), hosts



5.2 networks

NAME

networks - network name data base

DESCRIPTION

The networks file contains information regarding the known networks which comprise the DARPA Internet. For each network a single line should be present with the following information:

- official network name
- network number
- aliases

Items are separated by any number of blanks and/or tab characters. A ``#" indicates the beginning of a comment; characters up to the end of the line are not interpreted by routines which search the file. This file is normally created from the official network data base maintained at the Network Information Control Center (NIC), though local changes may be required to bring it up to date regarding unofficial aliases and/or unknown networks.

Network number may be specified in the conventional ``." notation using the `inet_network()` routine from the Internet address manipulation library, `inet()`. Network names may contain any printable character other than a field delimiter, newline, or comment character.

EXAMPLE


```
#
# Network entries
#
arpanet          10                arpa
matnet           128.14            # Mobile Access Terminal Net
aeronet          192.5.9           # Aerospace Labnet
olymp            192.1.1           # our own network
```

FILES

/DD/ETC/networks

SEE ALSO

`ifconfig()`, `getnetbyaddr()`, `getnetbyname()`, `setnetent()`, `getnetent()`, `endnetent()`



5.2.1 The struct netent structure

NAME

struct netent – the net entry structure


DESCRIPTION

The struct netent structure is used by the functions concerned with the network database and contains the broken-out fields of a line in the network data base, /DD/ETC/networks. it is defined as follows:

```
struct netent {  
    char          *n_name;          /* official name of net */  
    char          **n_aliases;      /* alias list */  
    int           n_addrtype;       /* net number type */  
    unsigned long n_net;            /* net number */  
};
```

The members of this structure are:

n_name	The official name of the network.
n_aliases	A zero terminated list of alternate names for the network.
n_addrtype	The type of the network number returned; currently only AF_INET.
n_net	The network number. Network numbers are returned in machine byte order.



5.2.2 getnetbyaddr()

NAME

getnetbyaddr() - get network entry

SYNOPSIS

```
#include <natdefs.h>
#include <netdb.h>
```

```
struct netent *getnetbyaddr(net, type)
    long net;
    int type;
```

DESCRIPTION

Getnetbyaddr() takes the parameters:

net	for the network number of the desired network
type	for the address type (currently only AF_INET)

RESULTS


If getnetbyaddr() has been executed successfully, it returns a pointer to a struct netent structure (see section 5.2.1, "The struct netent structure"). Otherwise, an error return status from getnetbyaddr() is indicated by return of a null pointer.

SEE ALSO

getnetbyname(), setnetent(), getnetent(), endnetent(), networks

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet address format is currently understood.



8.2.2 getnetbyname()

NAME

getnetbyname() - get network entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct netent *getnetbyname(name)
    char *name;
```

DESCRIPTION

Getnetbyname() takes the parameter:

name	for the name of the desired network
-------------	-------------------------------------

RESULTS


If getnetbyname() has been executed successfully, it returns a pointer to a struct netent structure (see section 5.2.1, "The struct netent structure"). Otherwise, an error return status from getnetbyname() is indicated by return of a null pointer.

SEE ALSO

getnetbyaddr(), setnetent(), getnetent(), endnetent(), networks

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet address format is currently understood.



5.2.4 getnetent()

NAME

getnetent() - get network entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct netent *getnetent()
```

DESCRIPTION

getnetent() always reads the next entry (line) of the networks data base, opening the file if necessary.

RESULTS


If getnetent() has been successfully executed, it returns a pointer to a struct netent structure (see section 5.2.1, "The struct netent structure"). Otherwise, an error return status from getnetent is indicated by return of a null pointer.

SEE ALSO

getnetbyaddr(), getnetbyname(), setnetent(), endnetent(), networks

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet address format is currently understood.



8.2.4 setnetent()

NAME

setnetent() - Open the networks data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>

setnetent(stayopen)
    int stayopen;
```


DESCRIPTION

The function setnetent() opens the networks data base and positions an internal pointer to the first entry in this data base. A subsequent call to getnetent() will return a pointer to the netent structure for the first entry in this data base (see section 5.2.1, "The struct netent structure").

If the parameter stayopen contains a non-zero value, the networks data base will remain open following subsequent calls to getnetbyaddr() or getnetbyname(). Normally, the networks data base is closed again following every function call.

SEE ALSO

getnetbyaddr(), getnetbyname(), getnetent(), endnetent(), networks



5.2.6 endnetent()

NAME

endnetent() - Close the networks data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>


endnetent()
```

DESCRIPTION

The function endnetent() closes the networks data base file after use.

SEE ALSO

getnetaddr(), getnetbyname(), setnetent(), getnetent(), networks



8.3 protocols

NAME

protocols - protocol name data base

DESCRIPTION

The protocols file contains information regarding the known protocols used in the DARPA Internet. For each protocol a single line should be present with the following information:

- official protocol name
- protocol number
- aliases

Items are separated by any number of blanks and/or tab characters. A ``#" indicates the beginning of a comment; characters up to the end of the line are not interpreted by routines which search the file.

Protocol names may contain any printable character other than a field delimiter, newline, or comment character.

EXAMPLE


```
#
# Internet (IP) Protocols
#
ip          0      IP      # Internet Protocol
icmp        1      ICMP    # Internet Control Message Protocol
ggp         3      GGP     # Gateway-Gateway Protocol
tcp         6      TCP     # Transmission Control Protocol
egp         8      EGP     # Exterior Gateway Protocol
pup         12     PUP     # PARC Universal Packet Protocol
udp         17     UDP     # User Datagram Protocol
hmp         20     HMP     # Host Monitoring Protocol
xns-idp     22     XNS-IDP  # Xerox NS IDP
rdp         27     RDP     # "Reliable Datagram" Protocol
```

FILES

/DD/ETC/protocols

SEE ALSO

getprotobyname(), getprotobynumber(), getprotoent(), setprotoent(), endprotoent()



5.3.1 The struct protoent structure

NAME

struct protoent – the protocol entry structure


DESCRIPTION

The struct protoent structure is used by the functions concerned with the network protocol data base and contains the broken-out fields of a line in the network protocol data base, /DD/ETC/protocols. It is defined as follows:

```
struct protoent {  
    char *p_name; /* official name of protocol */  
    char **p_aliases; /* alias list */  
    int p_proto; /* protocol number */  
};
```

The members of this structure are:

p_name	The official name of the protocol.
p_aliases	A zero terminated list of alternate names for the protocol.
p_proto	The protocol number.



5.3.2 getprotobyname()

NAME

getprotobyname() - get protocol entry

SYNOPSIS

```
#include <netdb.h>
```

```
struct protoent *getprotobyname(name)
    char *name;
```

DESCRIPTION

getprotobyname() takes the protocol name of the desired protocol in the parameter name.

RESULTS


If getprotobyname() has been successfully executed, it returns a pointer to a struct protoent structure (see section 5.3.1, "The struct protoent structure"). Otherwise, an error return status from getprotobyname() is indicated by return of a null pointer.

SEE ALSO

getprotobyname(), getprotoent(), setprotoent(), endprotoent(), protocols

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.3.3 getprotobynumber()

NAME

getprotobynumber() - get protocol entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct protoent *getprotobynumber(proto)
    int proto;
```

DESCRIPTION

getprotobynumber() takes the protocol number of the desired protocol in the parameter proto.

RESULTS


If getprotobynumber() has been successfully executed, it returns a pointer to a struct protoent structure (see section 5.3.1, "The struct protoent structure"). Otherwise, an error return status from getprotobynumber() is indicated by return of a null pointer.

SEE ALSO

getprotobyname(), getprotoent(), setprotoent(), endprotoent(), protocols

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.3.4 getprotoent()

NAME

getprotoent() - get protocol entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct protoent *getprotoent()
```

DESCRIPTION

The function getprotoent() reads the next line of the protocols data base, opening the file if necessary.

RESULTS


If getprotoent() has been successfully executed, it returns a pointer to a protoent structure (see section 5.3.1, "The struct protoent structure"). Otherwise, an error return status from getprotoent() is indicated by return of a null pointer.

SEE ALSO

getprotobyname(), getprotobynumber(), setprotoent(), endprotoent(), protocols

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.3.5 setprotoent()

NAME

setprotoent() - Open the protocols data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>

setprotoent(stayopen)
    int stayopen
```


DESCRIPTION

The function setprotoent() opens the protocols data base and positions an internal pointer to the first entry in this data base. A subsequent call to getprotoent() will return a pointer to the protoent structure for the first entry in this data base (see section 5.3.1, "The struct protoent structure").

If the parameter stayopen contains a non-zero value, the protocols data base will remain open following subsequent calls to getprotobyaddr() or getprotobyname(). Normally, the protocols data base is closed again following every function call.

SEE ALSO

getprotobyname(), getprotobynumber(), getprotoent(), endprotoent(), protocols



5.3.6 endprotoent()

NAME

endprotoent() - Close protocols data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>


endprotoent()
```

DESCRIPTION

The function endprotoent() closes the network protocol data base, /DD/ETC/protocols.

SEE ALSO

getprotobyname(), getprotobynumber(), getprotoent(), setprotoent(), protocols



5.4 services

NAME

services - services name data base

DESCRIPTION

The services file contains information regarding the known services available in the DARPA Internet. For each service a single line should be present with the following information:


- official service name
- port number
- protocol name
- aliases

Items are separated by any number of blanks and/or tab characters. The port number and protocol name are considered a single item; a ``/' is used to separate the port and protocol (e.g. ``512/tcp"). A ``#" indicates the beginning of a comment; subsequent characters up to the end of the line are not interpreted by the routines which search the file.

Service names may contain any printable character other than a field delimiter, newline, or comment character.

EXAMPLE

```
#
# Internet Network Services
#
echo          7/tcp
echo          7/udp
discard       9/tcp      sink null
discard       9/udp      sink null
daytime       13/tcp
daytime       13/udp
netstat       15/tcp
chargen       19/tcp      ttytst source
chargen       19/udp      ttytst source
ftp           21/tcp
telnet        23/tcp
smtp          25/tcpmail
```




FILES

/DD/ETC/services

SEE ALSO

getservbyname(), getservbyport(), setservent(), getservent(), endservent()



5.4.1 The struct servent structure

NAME

struct servent – the service entry structure


DESCRIPTION

The struct servent structure is used by the functions concerned with the network services data base and contains the broken-out fields of a line in the network services data base, /DD/ETC/services. It is defined as follows:

```
struct servent {  
    char *s_name; /* official name of service */  
    char **s_aliases; /* alias list */  
    int s_port; /* port service resides at */  
    char *s_proto; /* protocol to use */  
};
```

The members of this structure are:

s_name	The official name of the service.
s_aliases	A zero terminated list of alternate names for the service.
s_port	The port number at which the service resides. Port numbers are returned in network byte order.
s_proto	The name of the protocol to use when contacting the service.



5.4.2 getservbyname()

NAME

getservbyname() - get service entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct servent *getservbyname(name, proto)
    char *name, *proto;
```

DESCRIPTION

The function getservbyname() takes the port number of the desired service in the parameter port. A protocol name may be included for the service, however if it is it must agree with the port's entry in the services data base.

RESULTS


If getservbyname() has been successfully executed, it returns a pointer to a struct servent structure (see section 5.4.1, "The struct servent structure"). Otherwise, an error return status from getservbyname() is indicated by return of a null pointer.

SEE ALSO

getservbyport(), setservent(), getservent(), endservent(), services

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.4.3 getservbyport()

NAME

getservbyport() - get service entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct servent *getservbyport(port, proto)
    int port; char *proto;
```

DESCRIPTION

The function getservbyport() takes the port number of the desired service in the parameter port. A protocol name may be included for the service, however if it is it must agree with the port's entry in the services data base.

RESULTS


If getservbyport() has been successfully executed, it returns a pointer to a struct servent structure (see section 5.4.1, "The struct servent structure"). Otherwise, an error return status from getservbyport() is indicated by return of a null pointer.

SEE ALSO

getservbyname(), setservent(), getservent(), endservent(), services

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.4.4 getservent()

NAME

getservent() - get service entry

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>
```

```
struct servent *getservent()
```

DESCRIPTION

The function getservent() reads the next entry/line of the services data base, opening the file if necessary.

RESULTS


If getservent() has been successfully executed, it returns a pointer to a struct servent structure (see section 5.4.1, "The struct servent structure"). Otherwise, an error return status from getservent() is indicated by return of a null pointer.

SEE ALSO

getservbyname(), getservbyport(), setservent(), endservent(), services

NOTE

All information is contained in a static area so it must be copied if it is to be saved. Only the Internet protocols are currently understood.



5.4.5 setservent()

NAME

setservent() - Open services data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>

setservent(stayopen)
    int stayopen
```


DESCRIPTION

The function setservent() opens the services data base and positions an internal pointer to the first entry in this data base. A subsequent call to getservent() will return a pointer to the struct servent structure for the first entry in this data base (see section 5.4.1, "The struct servent structure").

If the parameter stayopen contains a non-zero value, the services data base will remain open following subsequent calls to getservbyaddr() or getservbyname(). Normally, the services data base is closed again following every function call.

SEE ALSO

getservbyname(), getservbyport(), getservent(), endservent(), services



5.4.6 endservent()

NAME

endservent() - Close the services data base

SYNOPSIS

```
#include <netdb.h>
#include <natdefs.h>

endservent()
```

DESCRIPTION

The function endservent closes the network services data base, /DD/ETC/services.

SEE ALSO

getservbyname(), getservbyport(), setservent(), getservent(), services